

Project Explanation: Online Payment Fraud Detection Model

This document explains the entire Python script used to build, train, and evaluate a machine learning model for detecting fraudulent online payments.

I. The Goal of the Code

The main goal is to create a highly accurate **Ensemble Model** (a "super-model" built from several smaller models) that can look at a new transaction and tell us if it is legitimate or **fraudulent** (the target we are trying to predict).

II. Understanding the Data Columns

Your model uses a dataset where each row is a transaction. Here are the most important columns and what they represent:

Column Name	Type	Purpose in Fraud Detection
isFraud	Integer (0 or 1)	Target Variable. This is what we want to predict: 1 means the transaction was fraud, and 0 means it was safe.
step	Integer	Represents the hour of the simulation. (e.g., step 1 is hour 1, step 743 is hour 743).
type	Text (Object)	The type of transaction, like CASH_OUT , TRANSFER , PAYMENT , or DEBIT .
amount	Decimal	The monetary amount of the transaction.
oldbalanceO rg	Decimal	The starting balance of the customer who made the transfer (nameOrig).
newbalanceO rig	Decimal	The balance of the customer after the transaction.
oldbalanceD est	Decimal	The starting balance of the recipient (nameDest).
newbalanceD est	Decimal	The balance of the recipient after the transaction.
isFlaggedFraud	Integer (0 or 1)	Whether the system flagged the transaction as suspicious based on simple internal rules (usually used only for analysis, not prediction).
nameOrig , nameDest	Text (Object)	Customer/Recipient IDs. These are typically dropped because they are identifiers that don't help the model generalize.

III. Step-by-Step Code Explanation

A. Setup and Data Loading

```

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline

data = pd.read_csv('...')
data.head()

```

- **Imports:** We bring in libraries (collections of tools):
 - `numpy` and `pandas` (`pd`): Essential for handling and manipulating data in tables (`DataFrames`).
 - `matplotlib.pyplot` (`plt`) and `seaborn` (`sns`): Used for making charts and graphs (visualizations).
 - `%matplotlib inline`: A special command to make charts appear right inside the notebook.
- **Data Loading:** The script reads your dataset (`new_data.csv`) into a pandas DataFrame named `data`. `data.head()` shows you the first few rows.

B. Initial Data Exploration (Data Types)

```

obj = (data.dtypes == 'object')
object_cols = list(obj[obj].index)
print("Categorical variables:", len(object_cols))
# ... (Similar code for 'int' and 'float' types)

```

- This section checks the type of data in each column (e.g., numbers, text, etc.).
- '**object**' type variables are usually text (like `type`, `nameOrig`, `nameDest`), which we call **Categorical Variables** (variables that represent categories, not measurable quantities). Machine learning models cannot directly process text, so they need special handling.

C. Visualizing Data Patterns (Exploratory Data Analysis)

```

sns.barplot(x='type', y='amount', data=data)
data['isFraud'].value_counts()

```

- **Bar Plot:** `sns.barplot(x='type', y='amount', data=data)` creates a bar chart comparing the average transaction amount across different transaction type categories. This helps see if certain transaction types involve higher amounts than others.
- **Value Counts:** `data['isFraud'].value_counts()` tells us exactly how many fraud (1) versus non-fraud (0) transactions are in the data. This is crucial for checking **Class Imbalance** (when one category is much rarer than others, which is very common in fraud data).

D. Correlation Heatmap (Feature Relationships)

```
sns.heatmap(data.apply(lambda x: pd.factorize(x)[0]).corr(),
             # ... plotting arguments
             annot=True)
```

- **Purpose:** This plot helps you quickly see how strongly every pair of columns relates to every other column.
- `pd.factorize(x)[0]` : This is a trick used to temporarily turn text columns (like `type`) into numbers so they can be included in the correlation calculation.
- **Heatmap:** A visual matrix where:
 - **Dark Color (Green/Teal):** Strong positive relationship (as one value goes up, the other tends to go up).
 - **Dark Color (Brown/Yellow):** Strong negative relationship (as one value goes up, the other tends to go down).
 - **Light Color (White):** Weak or no relationship.

E. Data Preparation for Modeling

```
type_new = pd.get_dummies(data['type'], drop_first=True) # ...
data_new = pd.concat([data, type_new], axis=1) # ...
X = data_new.drop(['isFraud', 'type', 'nameOrig', 'nameDest'], axis=1) # ...
y = data_new['isFraud'] # ...
```

- **One-Hot Encoding (`pd.get_dummies`):** This is how we handle the categorical `type` column. We convert the single 'type' column into several new number columns (e.g., `CASH_OUT`, `TRANSFER`, etc.). If a transaction is a 'TRANSFER', the `TRANSFER` column gets a `1`, and all others get a `0`.
- **Feature Selection:**
 - **x (Features):** This is the input data used to train the model. We drop columns that are not useful or are the target itself: `isFraud`, the old `type` column, and the identifier columns (`nameOrig`, `nameDest`).
 - **y (Target):** This is the output we want to predict (`isFraud`).

IV. Training the Machine Learning Models

A. Data Splitting

```
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.3, random_state=42)
```

- **Train-Test Split:** We divide our data into two parts:
 - **Training Data (70%):** Used to teach the model.

- **Test/Validation Data (30%):** Used to see how well the model works on data it has **never seen before**. This is the true test of performance.
- `random_state=42` : Ensures that every time you run the code, the split is the exact same, making your results **reproducible** (reliable and consistent).

B. Defining the Models

We define three different Machine Learning models, each specialized in finding patterns:

1. **Logistic Regression (`log_clf`):** A fundamental, simple model that finds a **linear relationship** between inputs and the probability of fraud.
2. **Random Forest Classifier (`rf_clf`):** An **ensemble model** (built from multiple simple decision trees) that is excellent at capturing complex, non-linear relationships.
 - `n_estimators=7` : We are using 7 individual decision trees.
3. **XGBoost Classifier (`xgb_clf`):** A highly efficient and advanced **ensemble model** that uses a technique called *Gradient Boosting*. It performed exceptionally well after fixing a critical issue:
 - `scale_pos_weight=773` (**CRITICAL FIX**): This tells the model to give the rare "Fraud" cases 773 times more importance than the common "Not Fraud" cases, directly solving the **Class Imbalance** problem.

C. Creating the Ensemble (The "Super-Model")

```
voting_clf = VotingClassifier(
    estimators=ensemble_models,
    voting='soft',
    n_jobs=-1
)
```

- **Voting Classifier:** This combines the predictions of the three models (`lr` , `xgb` , `rf`) to create a single, more robust prediction.
- `voting='soft'` : Means the final prediction is based on the **average probability** given by each model, which is best when evaluating models using the ROC AUC score.
- `n_jobs=-1` : Tells the script to use all available processing power (CPU cores) to train the models faster.

V. Training and Evaluation Metrics

A. Training the Final Model

```
voting_clf.fit(X_train, y_train)
```

- This is the moment of training! The `VotingClassifier` trains all three models simultaneously using the `X_train` data.

B. Final Evaluation Metrics

```
train_preds_ensemble = voting_clf.predict_proba(X_train)[:, 1]
print('Ensemble Training ROC AUC:', roc(y_train, train_preds_ensemble))
# ... (Validation ROC AUC is calculated similarly using X_test)
```